

Holo Fuel Model

Perry Kundert

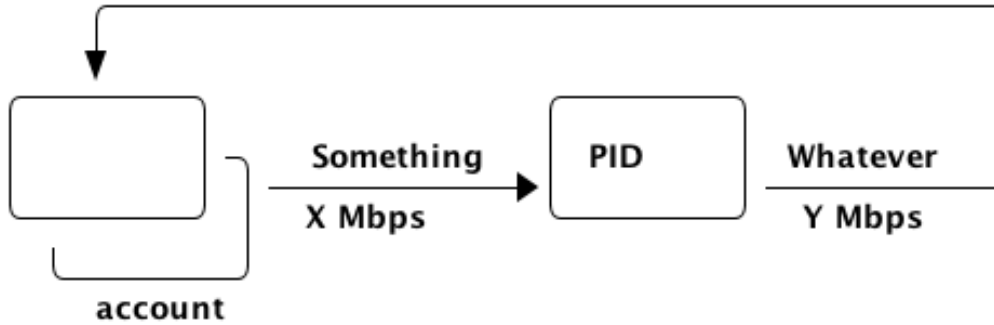
July 27, 2018

Contents

1	Holo Fuel Valuation	1
1.1	The Computational Resources Basket	2
1.1.1	Holo Hosting Premium	2
1.1.2	Resource Price Stability	3
1.2	Commodity Price Discovery	3
1.2.1	Recovering Commodity Basket Costs	3
1.2.2	Holo Hosting Receipts	5
1.2.3	Recovery of Commodity Valuations	6
1.2.4	Commodity Basket Valuation	6
1.3	Simple Value Stability Control via PID	7
1.3.1	Host/dApp Pricing	7
1.3.2	Host/dApp Pricing Automation Approaches	8
1.3.3	Simple Host/dApp Pricing Model	8
2	Holo Fuel Value Stabilization	11
2.1	Wealth Monetization	11
2.1.1	Reserve Accounts	11

1 Holo Fuel Valuation

A value-stable wealth-backed cryptocurrency platform, where each unit is defined in terms of a basket of computational resources, operating in a powerful decentralized verification environment.



1.1 The Computational Resources Basket

One Holo Fuel (HOT) is defined as being able to purchase 1 month of Holo Hosting services for the front-end (ie. web API, databases, etc.) portion of a typical dApp.

This might be roughly equivalent to the 2018 price of (and actual utilization of) a small cloud hosting setup (eg. several \$5/month Droplet on Digital Ocean at partial utilization hosting front-ends, DBs, backups, etc.), plus ancillary hosting services (represented by a premium for inclusion in the Holo system).

If the minimal Hosting costs for a small Web app is estimated at USD\$100/mo., and comprises 5 cloud hosting nodes and ancillary services for the various aspects of the system. An equivalent Holo Host based system would have similar CPU and storage requirements overall, but a greater redundancy (say, 5 x, so all DHTs spread across 25 Holo Hosts).

Thus, the basket of commodities defining the value of USD\$100 ~ 100 Holo Fuel could be defined as:

Commodity	Amount	Units	Weight	Description
holo	25.00	Host	60.000%	Inclusion in the Holo system
net	0.50	TB	20.000%	Internet bandwidth
ram	1.00	GB	5.000%	Processor memory
data	0.25	TB	10.000%	Persistent storage (DHT/DB/file)
cpu	1.00	Core	5.000%	A processing core

1.1.1 Holo Hosting Premium

A Holochain Distributed Application (dApp) hosted on Holo provides a valuable set of features, over and above simply hosting a typical web application on a set of cloud servers. These services must usually be either purchased, or architected by hand and distributed across multiple cloud hosting nodes for redundancy.

- ☐ Reliability. Few single points of failure.
- ☐ Backup. All DHT data is spread across many nodes.
- ☐ Scalability. Automatically scales to absorb increased load.

The value of Holo is substantial in terms of real costs to traditional app developers, and is a component of the basket of commodities defining the price of Holo Fuel. However, it's real

monetary value will emerge over time, as the developer community comprehends it. Our pricing algorithm must be able to dig this Holo premium component out of the historical hosting prices, as a separate component.

1.1.2 Resource Price Stability

There are many detailed requirements for each of these commodities, which may be required for certain Holochain applications; CPU flags (eg. AVX-512, cache size, ...), RAM (GB/s bandwidth), HDD (time to first byte, random/sequential I/O bandwidth), Internet (bandwidth/latency to various Internet backbone routers).

The relative distribution of these features will change over time; RAM becomes faster, CPU cores more powerful. The definition of a typical unit of these commodities therefore changes; as Moore's law decreases the price, the specifications of the typical computer also improve, counterbalancing this inflationary trend.

For each metric, the price of service on the median Holo Host node will be used; 1/2 will be below (weaker, priced at a discount), 1/2 above (more powerful, priced at a premium). This will nullify the natural inflationary nature of Holo Fuel, if we simply defined it in terms of fixed 2018 computational resources.

1.2 Commodity Price Discovery

Value stabilization requires knowledge of the current prices of each commodity in the currency's valuation basket, ideally denominated in the currency itself. If these commodities are traded within the cryptocurrency implementation, then we can directly discover them on a distributed basis. If outside commodity prices are used, then each independent actor computing the control loop must either reach consensus on the price history (as collected from external sources, such as Distributed Oracles), or trust a separate module to do so. In Holo Fuel, we host the sale of Holo Host services to dApp owners, so we know the historical prices.

When a history of Holo Hosting service prices is available, Linear Regression can be used to discover the average fixed (Holo Hosting premium) and variable (CPU, ...) component costs included in the prices, and therefore the current commodity basket price.

1.2.1 Recovering Commodity Basket Costs

To illustrate price recovery, let's begin with simulated prices of a basket of commodities. A prototypical minimal dApp owner could select 100 Holo Fuel worth of these resources, eg. 25x Holo Hosts, .05 TB data, 1.5 cpu, etc. as appropriate for their specific application's needs.

This Hosting selection wouldn't actually be a manual procedure; testing would indicate the kind of loads to expect for a given amount and type of user activity, and a calculator would estimate the various resource utilization and costs. At run time, the credit extended to the dApp owner (calculated from prior history of Hosting receipt payments) would set the maximum outstanding Hosting receipts allowed; the dApp deployment would auto-scale out to qualified Hosts in various tranches as required; candidate Hosts (hoping to generate Hosting receipts) would auto-install the application as it reached its limits of various resource utilization metrics across its current stable of Hosts.

```
def rnd_std_dst( sigma, mean=0, minimum=None ):
    """ """
    val = sigma * np.random.randn() + mean
```

```

    return val if minimum is None else max( minimum, val )

# To simulate initial pricing, lets start with an estimate of proportion of basket value represented
# by each amount of the basket's commodities. Prices of each of these commodities is free to float
# in a real market, but we'll start with some pre-determined "weights"; indicating that the amount
# of the specified commodity holds a greater or lesser proportion of the basket's value.
# Regardless, 100 Holo Fuel is guaranteed to buy the entire basket.
prices = {}
for k in basket:
    price_mean = basket_target * weight[k] / basket[k] # target price: 1 Holo Fuel == 1 basket / basket_target
    price_sigma = price_mean / 10 # difference allowed; about +/- 10% of target
    prices[k] = rnd_std_dst( price_sigma, price_mean )

[ [ "Commodity", "Price", "Per", "Per" ],
  None ] \
+ [ [ k, "%5.2f" % ( prices[k] ), commodities[k].units, 'mo.' ]
    for k in basket ]

```

Commodity	Price	Per	Per
holo	2.20	Host	mo.
net	34.48	TB	mo.
ram	5.61	GB	mo.
data	46.73	TB	mo.
cpu	4.76	Core	mo.

From this set of current assumed commodity prices, we can compute the current price of the Holo Fuel currency's basket:

```

basket_price = sum( basket[k] * prices[k] for k in basket )
[ [ "Holo Fuel Basket Price" ],
  None,
  [ "$%5.2f / %.2f" % ( basket_price, basket_target ) ] ]

```

$$\frac{\text{Holo Fuel Basket Price}}{\$94.30 / 100.00}$$


If the current price of this basket is >100 , then we are experiencing commodity price inflation; if <100 , price deflation. Feedback control loops will act to bring the price back to 100 Holo Fuel per basket.

```

labels = [ k for k in basket ]
sizes = [ basket[k] * prices[k] for k in basket ]
explode = [ .1 if k == 'holo' else 0 for k in basket ]
# with plt.xkcd():
fig1,ax1 = plt.subplots()
ax1.pie( sizes, explode=explode, labels=labels, autopct='%1.1f%%', shadow=True, startangle=90 )
ax1.axis( 'equal' ) # Equal aspect ratio ensures that pie is drawn as a circle.
plt.title( "%6.2f Holo Fuel Basket Price: %6.2f: %sflation" % (
    basket_target, basket_price, "in" if basket_price > basket_target else "de" ))
plt.show()

```

100.00 Holo Fuel Basket Price: 94.30: deflation



1.2.2 Holo Hosting Receipts

Once we have the currency's underlying commodity basket, let's simulate a sequence of trades of various amounts of these commodities. In the Holo system, this is represented by Hosts issuing receipts for services to dApp owners.

Each Hosting receipt will be for a single Holo Host, not for the entire dApp; the sum of all Holo Hosting receipts issued to the dApp owner for our archetypical small dApp would sum to approximately 100 Holo Fuel per month.

We will not know the exact costs of each commodity used to compute the price, or how much is the baseline Holo system premium. However, it will be dependant on the capability of the Host (stronger hosts can charge more, for hosting more specialized dApps), and the amount of various services used.

So, let's issue a bunch of small Holo Hosting receipts, each for approximately 1/25th of the total Holo Hosting load (since our small dApp is spread across 25 Holo Hosts).

```

amounts_mean      = 1.00
amounts_sigma     = 0.5
error_sigma       = 0.10 # +/- 10% variance in bids (error) vs. price
trades            = []
number            = 10000
for _ in range( number ):
    # Each dApp consumes a random standard distribution of the target amount of each commodity
    amounts        = { k: 1 if k == 'holo'
                       else basket[k] * rnd_std_dst( amounts_sigma, amounts_mean, minimum=0 ) / basket['holo']
                       for k in basket }
    price          = sum( amounts[k] * prices[k] for k in amounts )
    error          = price * rnd_std_dst( error_sigma )
    bid            = price + error
    trades.append( dict( bid = bid, price = price, error = error, amounts = amounts ))

[ [ "Fuel","calc/err", "dApp Requirements" ], None ] \
+ [ [
    "%5.2f" % t['bid'],
    "%5.2f%+5.2f" % ( t['price'], t['error'] ),
    ", ".join( "%5.4f %s %s" % ( v, k, commodities[k].units ) for k,v in t['amounts'].items() ),
    ]
  for t in trades[:5] ] \
+ [ [ '...' ] ]

```

Fuel	calc/err	dApp Requirements
3.52	3.29+0.23	1.0000 holo Host, 0.0114 net TB, 0.0031 ram GB, 0.0103 data TB, 0.0415 cpu Core
4.00	4.24-0.23	1.0000 holo Host, 0.0283 net TB, 0.0530 ram GB, 0.0138 data TB, 0.0251 cpu Core
3.56	4.04-0.48	1.0000 holo Host, 0.0275 net TB, 0.0241 ram GB, 0.0110 data TB, 0.0519 cpu Core
4.36	3.90+0.46	1.0000 holo Host, 0.0217 net TB, 0.0312 ram GB, 0.0125 data TB, 0.0411 cpu Core
3.99	4.05-0.07	1.0000 holo Host, 0.0141 net TB, 0.0644 ram GB, 0.0114 data TB, 0.0990 cpu Core
...		

1.2.3 Recovery of Commodity Valuations

Lets see if we can recover the approximate Holo baseline and per-commodity costs from a sequence of trades. Create some trades of 1 x Holo + random amounts of commodities around the requirements of a typical Holo dApp, adjusted by a random amount (ie. 'holo' always equals 1 unit, so that all non-varying remainder is ascribed to the "baseline" Holo Hosting premium).

Compute a linear regression over the trades, to try to recover an estimate of the prices.

```

items          = [ [ t['amounts'][k] for k in basket ] for t in trades ]
bids           = [ t['bid'] for t in trades ]

regression      = linear_model.LinearRegression( fit_intercept=False, normalize=False )
regression.fit( items, bids )
select         = { k: [ int( k == k2 ) for k2 in basket ] for k in basket }
predict        = { k: regression.predict( select[k] ) for k in basket }

[ [ "Score(R^2): ", "%.9r" % ( regression.score( items, bids ) ), ', ', '' ],
  None ] \
+ [ [ "Commodity", "Predicted", "Actual", "Error",
      # "selected"
    ],
  None ] \
+ [ [ k,
      "%.5.2f" % ( predict[k] ),
      "%.5.2f" % ( prices[k] ),
      "%.5.3f%%" % (( predict[k] - prices[k] ) * 100 / prices[k] ),
      #select[k]
    ]
  for k in basket ]

```

Score(R^2):	0.5623391		
Commodity	Predicted	Actual	Error
holo	2.20	2.20	+0.052%
net	34.39	34.48	-0.254%
ram	5.62	5.61	+0.218%
data	46.61	46.73	-0.252%
cpu	4.82	4.76	+1.276%

1.2.4 Commodity Basket Valuation

Finally, we can estimate the current Holo Fuel basket price from the recovered commodity prices.

```

basket_predict = sum( basket[k] * predict[k] for k in basket )
[ [ "Holo Fuel Price Recovered", "vs. Actual", "Error" ], None,
  [ "$%.2f / %.2f" % ( basket_predict, basket_target ),
    "%.5.2f" % ( basket_price ),
    "%.5.3f%%" % (( basket_predict - basket_price ) * 100 / basket_price ),
  ] ]

```

Holo Fuel Price Recovered	vs. Actual	Error
\$94.33 / 100.00	94.30	+0.030%

We have shown that we should be able to recover the underlying commodity prices, and hence the basket price with a high degree of certainty, even in the face of relatively large differences in the mix of prices paid for hosting.

1.3 Simple Value Stability Control via PID

The simplest implementation of value-stability is to directly control the credit supply. Lets establish a simple wealth-backed monetary system with a certain amount of wealth attached to it, from which we extend credit at a factor K of 0.5 to begin with; half of the value of the wealth is available in credit. Adjusting K increases/reduces the liquid credit supply.

The economy has a certain stock of Host resources available (eg. cpu, net, ...), and a certain pool of dApp owners wanting to buy various combinations of them. The owners willing to pay more will get preferred access to the resources. In a traditional bid/ask market, greater bids are satisfied first, lesser later or not at all. In Holo, tranches of similar Hosts round-robin requests from clients of the dApps they host.

1.3.1 Host/dApp Pricing

In the Holo Host environments, Hosts are pooled in tranches of like resource capacity (eg. cpu: type, count, ...), quality (eg. service: availability, longevity, ...), and price (eg. autopilot/manual pricing: lolo, lo, median, hi, hihi). A multi-dimensional table of Host tranches is maintained; each Host inserts itself into the correct table.

- TODO: How do the DHT peers confirm that a Host isn't lying about its internal computational resources? A dApp could check, and issue a warrant if the Host is lying, but a DHT peer couldn't independently verify these claims. There will be great incentive to inflate claims, to draw and serve higher-priced requests...)

A dApp owner also selects the resource requirements (eg. cpu: avx-128+, 4+ cores, ...) service level and pricing (eg. median).

Requests from hihi priced dApps are distributed first to the lolo, then lo, ..., hihi tranches of Hosts, as each tranche's resources is saturated; thus, lolo priced Hosts are saturated first. Then, hi dApps are served any by lolo, lo, ... Hosts not yet saturated, and so on. Thus, in times of low utilization (less dApps than Hosts), the highest priced Hosts may remain idle; in high utilization (more dApps than Hosts), the lowest priced dApp's requests may remain unserved (or, perhaps throttled and served round-robin, to avoid complete starvation of the lower priced dApp groups). Of course, these tranches of Hosts are also limited (via a set Union) to those Hosts in each tranche that **also** host a given target dApp, and requests for a dApp are only sent to those hosts who can service it.

- TODO: Each TCP/IP HTTP socket, representing 1 or more HTTP requests or a WebSocket initiation, is assigned a Host; does Holo terminate the connection and relay I/O to/from the Host? It should pre-establish a pool of sockets to candidate Hosts, ready to be distributed to incoming requests, thus eliminating the delay of the 3-way handshake, and pre-eliminating dead/unreachable Hosts.) This requires a persistent proxy a.la. Cloudflare. Much more simply, perhaps, we could build DNS servers that advertise multiple A records from an appropriate tranche of candidate servers, in round-robin fashion, and let the end-user sort out servers that disappear (until the DNS server figures out they're dead and stops serving their IP address). However, intervening caching DNS servers (eg. at large ISPs) could conduit large

numbers of request (ie. from the entire ISP!) to those few Host A-records for the time-to-live of the cached DNS query.

1.3.2 Host/dApp Pricing Automation Approaches

How does the system compute the actual price that "median" Hosts get paid? How does it evolve over time? 1/2 of requests should go to median, lo, lolu Hosts, and 1/2 should go to median, hi, hihi Hosts. A PID loop could move the "Median" Host price to make this true, perhaps. Hosts should set a minimum average price they'll earn, dApps a maximum average price they're willing to pay, and their requests are throttled to only the Host tranches which satisfy these limits.

By automatically switching a Host to higher/lower pricing tiers, and the dApp to lower/higher pricing selections, as their limit prices are reached, the numbers of Hosts/dApps above/below "median" changes – and the PID loop adjusts the median price to achieve above/below equilibrium. Thus, as more dApps exceed their high limit, switch to lower tiers (eg. from hi → median → lo), the mix of requests above/below median price changes, and the PID loop responds by adjusting the median Hosting price, which affects average dApp request pricing, which causes the dApp to hit its limits, which causes it to (again) switch to a lower tier...

Of course, the dApp owner is informed of this, in real time, and can make price limit adjustments, to re-establish dApp performance. Likewise, a Hosting owner can see that their Hosts are saturated/idle, and increase/decrease their minimum price, or maximum utilization targets; the Host should increase its desired pricing tier, to stay under its maximum utilization target.

1.3.3 Simple Host/dApp Pricing Model

For the purposes of this simple test, we'll assume that the Host will simply spend all the credit the dApp has available serving its requests (we won't simulate the dApps). So, let's generate a sequence of request service receipts from the Host to dApp owners, tuned to the credit available to the dApp.

```
# Adjust this so that our process value 'basket_value' achieves setpoint 'basket_target'
class K( object ): # Simplest, static K-valued
    value = 0.5

current_prices = dict( prices )

def current_value( basket_prices=None, currency_basket=None ):
    """Compute the value of a currency_basket (default: global basket) at some basket_prices
    (default: global prices)"""
    return sum( ( basket_prices or prices )[k] * ( currency_basket or basket )[k]
    for k in prices )

duration_hour = 60 * 60
duration_day = 24 * duration_hour
duration_month = 365.25 * duration_day / 12 # 2,629,800s.

used_mean      = 1.0 # Hourly usage is
used_sigma     = used_mean * .10 # +/-10%
reqs_mean      = 2.0      # Avg. Host is 2x minimal
reqs_sigma     = reqs_mean * .5 # +/- 50%

class dApp( object ):
    def __init__( self, duration=duration_month ): # 1 mo., in seconds
        """Select a random basket of computational requirements, ~ 2 +/- 1 times the minimal dApp,
        represented by the Holo Fuel basket, for the specified duration (min. 10% of basket). """
        self.duration = duration
        self.requires = { k: rnd_std_dst( sigma=1.0, mean=2.0, minimum=1/10 ) * basket[k] * duration / duration_month
        for k in basket }
    # Finally, compute the wealth required to fund this at current credit factor K
```



```

self.wealth = current_value( current_prices, self.requires ) / K.value

    def credit( self, dt=None ):
        """Credit available for dt seconds (1 hr., default) of Hosting."""
        return self.wealth * K.value * ( dt or duration_hour ) / self.duration

    def used( self, dt=None, mean=1.0, sigma=.1 ):
        """Resources used over period dt (+/- 10% default, but at least 0)"""
        return { k: self.requires[k] * rnd_std_dst( sigma=sigma, mean=mean, minimum=0 ) * dt / self.duration
                for k in self.requires }

class Host( object ):
    def __init__( self, dApp ):
        self.dApp = dApp

    def receipt( self, dt=None ):
        """Generate receipt for dt seconds worth of hosting our dApp; spend all available credit. This
        could be achieved, for example, by selecting a lower pricing tier (thus worse performance)."""
        avail = self.dApp.credit( dt=dt ) # Credit available
        used = self.dApp.used( dt=dt, sigma=used_sigma ) # Roughly what hosting resources would be used
        value = current_value( current_prices, used ) # total value of dApp Hosting resources used
        #print( "avail: {}, value: {}, used: {}".format( avail, value, used ) )
        return avail,used

hosts_count = 60 * 60 # ~1 Hosting receipt per second
hosts = [ Host( dApp() ) for _ in range( hosts_count ) ]
hours_count = 24

class K_sine( K ):
    sin_amp = K.value * 1/100 # +/- 1% of K
    sin_theta = 0
    sin_step = 2 * math.pi / hosts_count / 6 # 1 cycle per 6 hours
    def advance( self ):
        self.sin_theta += self.sin_step
    @property
    def value( self ):
        return super( K_sine, self ).value + self.sin_amp * math.sin( self.sin_theta )

K = K_sine()

receipts = []
for _ in range( hours_count ):
    for h in hosts:
        receipts.append( h.receipt( dt=duration_hour ) )
    K.advance()

cost,rcpt = receipts[-1]

items          = [ [ rcpt[k] for k in basket ] for cost,rcpt in receipts ]
costs          = [ cost for cost,rcpt in receipts ]

regression     = linear_model.LinearRegression( fit_intercept=False, normalize=False )
regression.fit( items, costs )
select         = { k: [ int( k == k2 ) for k2 in basket ] for k in basket }
predict        = { k: regression.predict( select[k] ) for k in basket }

actual_value = current_value( current_prices )
predict_value = current_value( predict )
[ [ "1hr Hosting Cost" ] + list( rcpt.keys() ),
  None,
  [ "%.6f" % cost ] + list( "%.6f" % rcpt[k] for k in rcpt ),
  None,
  [ "%dhr. x %d Hosts Cost" % ( hours_count, hosts_count ) ],
  None,
  [ "%.6f" % sum( cost for cost,rcpt in receipts ) ] + [ "%.6f" % sum( rcpt[k] for cost,rcpt in receipts ) for k in basket ],
  None,
  [ "Score(R^2) %.9r" % ( regression.score( items, costs ) ) ],
  [ "Predicted" ] + [ "%.2f" % predict[k] for k in basket ],

```

```

[ "Actual" ]      + [ "%5.2f" % current_prices[k] for k in basket ],
[ "Error" ]       + [ "%+5.3f%%" % (( predict[k] - current_prices[k] ) * 100 / current_prices[k] ) for k in basket ],
None,
[ "Actual Basket", "%5.2f" % actual_value ],
[ "Predict Basket", "%5.2f" % predict_value ],
[ "Error" , "%+5.3f%%" % (( predict_value - actual_value ) * 100 / actual_value ) ],
]

```

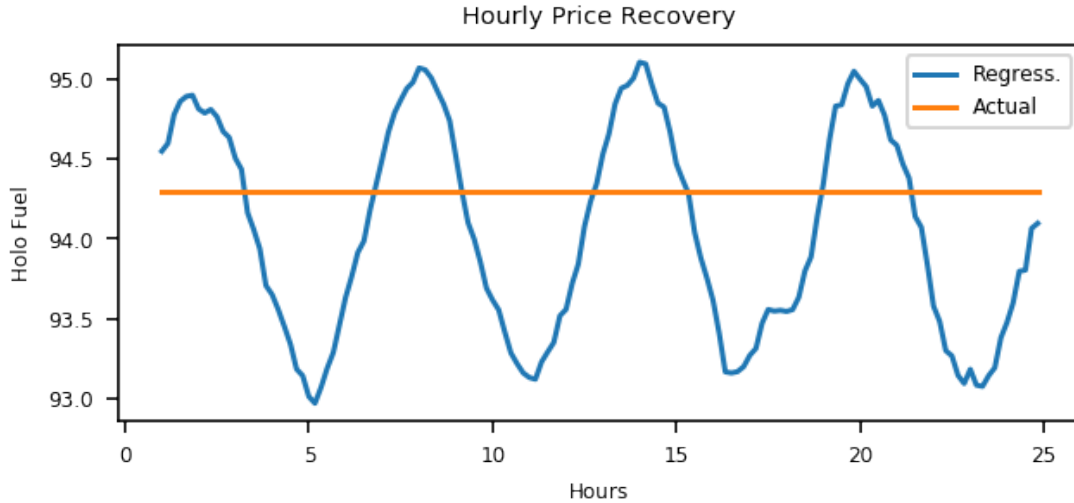
1hr Hosting Cost	holo	net	ram	data	cpu
0.265627	0.060328	0.001723	0.001556	0.000963	0.004739
24hr. x 3600 Hosts Cost					
22266.887759	5858.870686	120.119522	240.329732	59.127323	236.103865
Score(R ²) 0.9499280					
Predicted	2.13	34.53	6.23	47.68	5.43
Actual	2.20	34.48	5.61	46.73	4.76
Error	-3.269%	+0.136%	+11.043%	+2.023%	+14.096%
Actual Basket	94.30				
Predict Basket	94.05				
Error	-0.264%				

Lets see how well an hourly linear regression tracks the actual Basket price, in 10 minute intervals (so, 6 x 1-hour regression samples per hour). Lets see if we can pick up the 1% sine-wave variation in Credit Factor K every 6 hours:

```

# x is the fractional hour of the ends of each hour-long segment
x_divs = 6
x = [ s / x_divs + 1 for s in range( hours_count * x_divs ) ]
reg = []
act = []
for h in x: # Compute beg:end indices from fractional hour at end of each 1-hour range
    beg,end = int( (h-1) * hosts_count ),int( h * hosts_count )
    items      = [ [ rcpt[k] for k in basket ] for cost,rcpt in receipts[beg:end] ]
    costs      = [ cost for cost,rcpt in receipts[beg:end] ]
    regression.fit( items, costs )
    select     = { k: [ int( k == k2 ) for k2 in basket ] for k in basket }
    predict    = { k: regression.predict( select[k] ) for k in basket }
    reg.append( current_value( predict ))
    act.append( current_value( current_prices ))
plt.plot( x, reg, label="Regress." )
plt.plot( x, act, label="Actual" )
plt.xlabel( "Hours" )
plt.ylabel( "Holo Fuel" )
plt.legend( loc="upper right" )
plt.title( "Hourly Price Recovery" )
plt.show()

```



2 Holo Fuel Value Stabilization

Price discovery gives us the tools we need to detect {in,de}flation as it occurs. Control of liquid credit available in the marketplace gives us the levers we need to eliminate it.

Traditional Fiat currencies control the issuance of liquidity by influencing the commercial banks to create more or less money through lending, and to increase/reduce liquidity through the net issuance/retirement of debt (which creates/destroys the principal money).

2.1 Wealth Monetization

In a wealth-backed currency, credit is created by the attachment of wealth to the monetary system, and credit lines of varying proportions being extended against the value of that wealth.

Depending on savings rates, monetary velocity, public sentiment etc., the amount of credit available to actually be spent varies. Since this available liquid credit is split between possible expenditures in priority order, the amount available to spend on each specific commodity therefore varies, driving the market price up and down.

If reliable indicators of both the liquid credit supply within, and the quality and amount of wealth attached, exist within the system itself then control systems can be executed within the system to automatically control the monetization of wealth to achieve credit unit value equilibrium – value-stability.

Each reserve of wealth provided different flows and indicators, and can support value-stability in different ways.

2.1.1 Reserve Accounts

The Reserve Accounts provide the interface between external currencies (eg. USD\$, HOT ERC20 Tokens) and Holo Fuel.

Deposits to the reserve creates Holo Fuel credit limit (debt) at a current rate of exchange (TBD; eg. market rate + premium/discount). The corresponding Holo Fuel credits created are deposited to the recipient's account.

If Holo Fuel inflation occurs within the system, credit must be withdrawn. One way to accomplish this is to discourage creation of Holo Fuel (and encourage the redemption of Holo Fuel), by

increasing the exchange rate. The inverse (lowering exchange rate) would result in more Holo Fuel creation (less redemption), reducing the Holo Fuel available, and thus reduce deflation.

The Reserve Accounts can respond very quickly, inducing Holo Hosts with Holo Fuel balances to quickly convert them out to other currencies when exchange rates rise. Inversely, reducing rates would release waiting dApp owners to purchase more Holo Fuel for hosting their dApps, deploying it into the economy to address deflation (increasing computational commodity prices).

A PD (Proportional Differential) control might be ideal for this. This type of control responds quickly both to direct errors (things being the wrong price), but most importantly to changes in the 2nd derivative (changes in rate of rate of change); eg. things getting more/less expensive at an increasing rate.

By eliminating the I (Integral) component of the PID loop, it does **not** slowly build up a systematic output bias; it simply adjusts the instantaneous premium/discount added to the current market exchange rate (eg. the HOT ERC20 market), to arrive at the Reserve Account exchange rate. When inflation/deflation disappears, then the Reserve Account will have the same exchange rate as the market.

Beginning with a set of reserves:

```
reserve_t = collections.namedtuple(
    'Reserve', [
        'rate',      # Exchange rate used for these funds
        'amount',    # The total value of the amount executed at .rate
    ] )           # and the resultant credit in Holo Fuel == amount * rate

reserve = {
    'EUR':          [],          # LIFO stack of reserves available
    'USD':          [ reserve_t( .0004, 200 ), reserve_t( .0005, 250 ) ], # 1,000,000 Holo Fuel
    'HOT ERC20':    [ reserve_t( 1, 1000000 ) ], # 1,000,000 Holo Fuel
}

def reserves( reserve ):
    return [ [ "Currency", "Rate avg.", "Reserves", "Holo Fuel Credits", ], None, ] \
    + [ [ c, "%8.6f" % ( sum( r.amount * r.rate for r in reserve[c] )
        / ( sum( r.amount for r in reserve[c] ) or 1 ) ),
        "%8.2f" % sum( r.amount for r in reserve[c] ),
        "%8.2f" % sum( r.amount / r.rate for r in reserve[c] ) ]
        for c in reserve ] \
    + [ None,
        [ '', '', '', sum( sum( r.amount / r.rate for r in reserve[c] ) for c in reserve ) ] ]

summary = reserves( reserve )
summary # summary[-1][-1] is the total amount of reserves credit available, in Holo Fuel
```

Currency	Rate avg.	Reserves	Holo Fuel Credits
HOT ERC20	1.000000	1000000.00	1000000.00
USD	0.000456	450.00	1000000.00
EUR	0.000000	0.00	0.00
			2000000.0

As a simple proxy for price stability, lets assume that we strive to maintain a certain stock of Holo Fuel credits in the system for it to be at equilibrium. We'll randomly do exchanges of Holo Fuel out through exchanges at a randomly varying rate (also varied by the rate premium/discount), and purchases of Holo Fuel through exchanges at a rate proportional to the premium/discount.

```
t_last = -1
for t in range( 1000 ):
    dt = t - t_last
```